

# Accelerating inference of YOLOv3-tiny model using FPGA

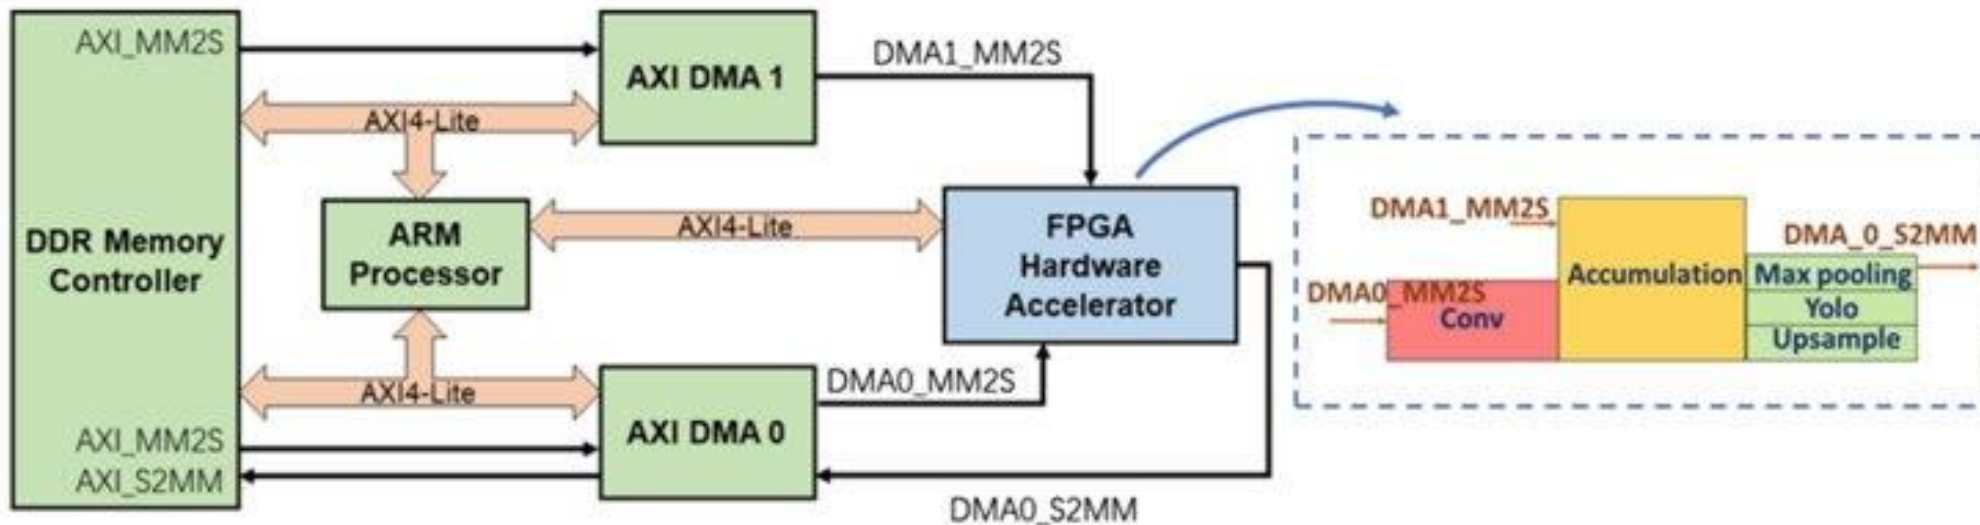
Anshuman  
Arvind Nataraj Sivasankar  
Shreyas Tater

# Problem

- Deploy YOLOv3-tiny on FPGA with minimal reduction in object detection precision
- Limited resources on FPGA (off chip memory required to store parameters and partial DNN inference computation results)
- Low Latency for real time object detection applications on embedded systems

# Design

- The FPGA accelerator consists of a three-stage pipeline.
- The first stage of the pipeline supports the execution of the convolution layer, whose output is accumulated in the second stage of the pipeline.
- Depending on the network structure that is executed at a given time, the accumulation results are sent for further processing in the Max pooling, Upsample or Yolo layer



# Approach

- Worked on a golden C code based off available literature - used \*pointers.
- Redid golden C code using Lab2 as base and made modifications to get single layer convolution running for a variable input size.

```
void tiled_conv_maxpool_id3 (  
    fm_t input_feature_map[1024][416][416],  
    wt_t layer_conv_weights[1024][1024][3][3],  
    fm_t output_feature_map[1024][416][416],  
    int id,  
    int ih,  
    int iw,  
    int od  
)
```

Array size matches the max dimensions that is possible

Actual layer input/output size bounds being passed

# Approach

Layer	Type	Filters	Size/Stride	Input	Output
0	Convolutional	16	$3 \times 3/1$	$416 \times 416 \times 3$	$416 \times 416 \times 16$
1	Maxpool		$2 \times 2/2$	$416 \times 416 \times 16$	$208 \times 208 \times 16$
2	Convolutional	32	$3 \times 3/1$	$208 \times 208 \times 16$	$208 \times 208 \times 32$
3	Maxpool		$2 \times 2/2$	$208 \times 208 \times 32$	$104 \times 104 \times 32$
4	Convolutional	64	$3 \times 3/1$	$104 \times 104 \times 32$	$104 \times 104 \times 64$
5	Maxpool		$2 \times 2/2$	$104 \times 104 \times 64$	$52 \times 52 \times 64$
6	Convolutional	128	$3 \times 3/1$	$52 \times 52 \times 64$	$52 \times 52 \times 128$
7	Maxpool		$2 \times 2/2$	$52 \times 52 \times 128$	$26 \times 26 \times 128$
8	Convolutional	256	$3 \times 3/1$	$26 \times 26 \times 128$	$26 \times 26 \times 256$
9	Maxpool		$2 \times 2/2$	$26 \times 26 \times 256$	$13 \times 13 \times 256$
10	Convolutional	512	$3 \times 3/1$	$13 \times 13 \times 256$	$13 \times 13 \times 512$
11	Maxpool		$2 \times 2/1$	$13 \times 13 \times 512$	$13 \times 13 \times 512$
12	Convolutional	1024	$3 \times 3/1$	$13 \times 13 \times 512$	$13 \times 13 \times 1024$
13	Convolutional	256	$1 \times 1/1$	$13 \times 13 \times 1024$	$13 \times 13 \times 256$
14	Convolutional	512	$3 \times 3/1$	$13 \times 13 \times 256$	$13 \times 13 \times 512$
15	Convolutional	255	$1 \times 1/1$	$13 \times 13 \times 512$	$13 \times 13 \times 255$
16	YOLO				
17	Route 13				
18	Convolutional	128	$1 \times 1/1$	$13 \times 13 \times 256$	$13 \times 13 \times 128$
19	Up-sampling		$2 \times 2/1$	$13 \times 13 \times 128$	$26 \times 26 \times 128$
20	Route 19 8				
21	Convolutional	256	$3 \times 3/1$	$13 \times 13 \times 384$	$13 \times 13 \times 256$
22	Convolutional	255	$1 \times 1/1$	$13 \times 13 \times 256$	$13 \times 13 \times 256$
23	YOLO				

- Golden – C
  - Implemented all layers – YOLO, Convolution, Max Pooling, Upsample and Route
  - Used Lab 2 code as base to develop code.
  - Created a 24-layer network

# Code Snippets

- YOLO C-Model code


```
for (int f = 0; f < 2; f++)
    for(int i = 0; i < (input_h); i++)
        for(int j = 0; j < (input_w); j++)
        {
            input_dxdy0[f][i][j] = (sigmoid(input_dxdy0[f][i][j]) + xy_grid0[f][i][j]) * stride;
            input_dxdy1[f][i][j] = (sigmoid(input_dxdy1[f][i][j]) + xy_grid1[f][i][j]) * stride;
            input_dxdy2[f][i][j] = (sigmoid(input_dxdy2[f][i][j]) + xy_grid2[f][i][j]) * stride;
            input_dwdh0[f][i][j] = fm_t(exp(input_dwdh0[f][i][j])) * *(anchor + f) * stride;
            input_dwdh1[f][i][j] = fm_t(exp(input_dwdh1[f][i][j])) * *(anchor + 2 + f) * stride;
            input_dwdh2[f][i][j] = fm_t(exp(input_dwdh2[f][i][j])) * *(anchor + 4 + f) * stride;
```

# HLS Code Approach

- Had to modify network to implement on FPGA
- Did not implement Batch Normalization on FPGA
- Faced implementation problems for sigmoid and exponential activation – did not implement the YOLO layer on FPGA
- Used Leaky ReLU for convolution


## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



## Leaky ReLU


$$\max(0.1x, x)$$



# Code Snippets

- Max Pooling

```
MAXPOOL_DEPTH:
for (int d = 0; d < 16; d++)
{
    MAXPOOL_HEIGHT:
    for(int h = 0; h < 208; h++)
    {
        MAXPOOL_WIDTH:
        for(int w = 0; w < 208; w++)
        {
            max = in_buf[d][2*h][2*w];
            MAXPOOL_WIN_DIM1:
            for (int i = 0; i < 2; i++)
            {
                MAXPOOL_WIN_DIM2:
                for(int j = 0; j < 2; j++)
                {
                    if (in_buf[d][(2*h)+i][(2*w)+j] > max)
                        max = in_buf[d][(2*h)+i][(2*w)+j];
                }
            }
            out_buf[d][h][w] = max;
        }
    }
}
```





# Optimization Strategies

- Array Partitioning
- Loop Pipelining
- Loop Unroll

```
#pragma HLS array_partition variable=X_buf dim=1
#pragma HLS array_partition variable=X_buf dim=2 factor=3 cyclic
#pragma HLS array_partition variable=X_buf dim=3 factor=3 cyclic
#pragma HLS array_partition variable=W_buf dim=2
#pragma HLS array_partition variable=W_buf dim=3
#pragma HLS array_partition variable=W_buf dim=4

OUT_DEPTH:
for (int outDepth = 0; outDepth < OUT_BUF_DEPTH; outDepth++)
{
    OUT_HEIGHT:
    for (int outHeight = 0; outHeight < OUT_BUF_HEIGHT; outHeight++)
    {
        OUT_WIDTH:
        for (int outWidth = 0; outWidth < OUT_BUF_WIDTH; outWidth++)
        {
#pragma HLS pipeline II=1
            Y_buf[outDepth][outHeight][outWidth] = 0;
            fm_t local[IN_BUF_DEPTH];
            IN_HEIGHT:
            for (int kHeight = 0; kHeight < 3; kHeight++)
            {
                IN_WIDTH:
                for (int kWidth = 0; kWidth < 3; kWidth++)
                {
```

# Optimization Strategies

- Loop Tiling
- Ping Pong Buffer

```
outD:
for (int outDepthOffset = 0; outDepthOffset < 16 / OUT_BUF_DEPTH; outDepthOffset++)
{
    // PING PONG
    // PING LOAD
    load_input_tile_block_from_DRAM(conv_in_buf, input_feature_map, ti, tj, 0);
    load_layer_params_from_DRAM(conv_wt_buf, conv_bias_buf, layer_weights, layer_bias, o
inD:
for (int inDepthOffset = 0; inDepthOffset < (3 / IN_BUF_DEPTH) - 1; inDepthOffset++)
{
    if (inDepthOffset % 2 == 0)
    {
        // PING COMPUTE + PONG LOAD
        conv_3x3(conv_out_buf, conv_in_buf, conv_wt_buf);
        save_partial_output_tile_block(partial_out_fm_buf, conv_out_buf, conv_bias_b
        load_input_tile_block_from_DRAM(conv_in_buf_1, input_feature_map, ti, tj, in
        load_layer_params_from_DRAM(conv_wt_buf_1, conv_bias_buf_1, layer_weights, L
    }
    else
    {
        // PONG COMPUTE + PING LOAD
        conv_3x3(conv_out_buf, conv_in_buf_1, conv_wt_buf_1);
        save_partial_output_tile_block(partial_out_fm_buf, conv_out_buf, conv_bias_b
        load_input_tile_block_from_DRAM(conv_in_buf, input_feature_map, ti, tj, inDe
        load_layer_params_from_DRAM(conv_wt_buf, conv_bias_buf, layer_weights, layer
```

# Evaluation

- Unoptimized, Synthesizable Run for all layers:

Latency (cycles)		Latency (absolute)		Interval		Pipeline
min	max	min	max	min	max	Type
367148894	47550818342	3.671 sec	475.508 sec	367148895	47550818343	none

- Accelerated Run:

Latency (cycles)		Latency (absolute)		Interval		Pipeline
min	max	min	max	min	max	Type
68493294	4029377062	0.685 sec	40.294 sec	68493295	4029377063	none

# Evaluations – Single layer data

	Latency (sec)	SpeedUp
Unoptimized	3.2	1
Ping Pong Buffer	2.02	1.584158
Loop Unroll + Pipelining	0.186	17.2043
Array Partitioning	0.0549	58.2878

# Evaluation

- Accelerated run for four layers:

```
=====
== Utilization Estimates
=====
* Summary:
```

Name	BRAM_18K	DSP	FF	LUT	URAM
DSP	-	4	-	-	-
Expression	-	-	0	820	-
FIFO	-	-	-	-	-
Instance	225	129	20332	41581	-
Memory	-	-	-	-	-
Multiplexer	-	-	-	1151	-
Register	-	-	1099	64	-
Total	225	133	21431	43616	0
Available	432	360	141120	70560	0
Utilization (%)	52	36	15	61	0

Latency:

\* Summary:

Latency (cycles)		Latency (absolute)	
min	max	min	max
30498837	30498837	0.305 sec	0.305 sec

# Future Actions to try before submitting code

- Our initial approach caused issues with debugging.
- Use independent functions for all convolution layers.
- Literature search for references with YOLO layer implementation on FPGA